

AI Pod Staffing Solution

Business and Functional Requirements



Date: July 17, 2026, Version 1.0

Project Name	AI Pod Staffing Solution	
Teams	<i>Customer Programs & OU</i>	<i>Innovation Studio</i>
Executive Sponsors	Amy Lawrence	Srinivas Thella
Project Leads	Indranie Balkaran	Sathish Chandrasekaran
Project Members	Robert Story, Brenna Cooper, Umu Rodon Rigau	TBC
Status	Draft for Review	Date: 7/17/26

1. Executive Summary

The AI Pod Staffing Solution introduces a standardized operating model for assembling cross-functional project teams, or Pods, using AI to assist with staffing recommendations. The solution is intended to improve the speed, consistency, transparency, and scalability of staffing decisions while preserving human accountability and governance. The proposed capability centralizes request intake, captures delivery and capability taxonomies, maintains experience and capacity profiles, and generates explainable AI recommendations for Pod Leads and Pod Members. Final assignments remain subject to human review and approval.

This document defines the business problem, desired future state, scope boundaries, process model, functional requirements, AI recommendation framework, reporting needs, governance, non-functional requirements, MVP scope, and success measures. It is intended to support executive alignment, business analysis, solution design, development, testing, and pilot execution.

2. Business Problem

- Requests arrive through multiple channels, which makes demand difficult to track consistently.
- There is no single repository for requests, assignments, and outcomes.
- Staffing decisions rely heavily on personal knowledge and informal judgment.
- The same resources may be repeatedly tapped while broader talent across the team remains underutilized.
- Visibility into capacity, interests, experience, and demand is limited.
- There is no consistent method for matching work to the most appropriate resources.
- Reporting on staffing effectiveness, portfolio demand, and delivery outcomes is fragmented or absent.
- These gaps make staffing slower, less transparent, and difficult to scale as demand grows.

3. Business Objectives

- Reduce the time required to staff requests.
- Improve the quality and consistency of resource allocation.
- Increase visibility into organizational capability, capacity, and demand.
- Use AI to recommend staffing options based on business data rather than ad hoc judgment alone.
- Preserve human governance, accountability, and approval of all assignments.
- Create a scalable operating model that can support pilot and broader adoption.

4. Scope (In / Out)

In Scope

- Standardized request intake.
- Centralized request repository.
- Delivery taxonomy and capability taxonomy.
- Experience profiles, interests, and capacity tracking.
- AI-assisted Pod Lead and Pod Member recommendations.
- Human review and approval workflow.
- Governance logging and audit trail.
- Operational and leadership dashboards.

Out of Scope for the MVP

- Automatic assignment without human approval.
- Solution is not meant to be a project management tool; rather it is limited to a Staffing Recommendation solution

5. Current-State Assessment

The current staffing process is informal and distributed. Requests may be submitted through email, Slack, meetings, or direct conversations. Because there is no single intake path or repository, it is difficult to track request status, understand demand patterns, or compare potential resource assignments objectively. Experience and availability are known by individuals rather than maintained in a structured and reusable form. As a result, the process depends on personal knowledge and does not scale well when demand increases.

Baseline measurements should be captured during analysis, including average staffing time, request volume, number of active requests, current utilization patterns, and the degree to which assignments are repeated across the same resources.

6. Future-State Operating Model

- People remain accountable for request submission, review, approval, and delivery.
- AI supports staffing decisions by analyzing request needs against structured business data.
- The solution standardizes request intake and creates a single source of truth for demand and assignment history.
- Governance ensures that every recommendation is reviewed before it is finalized.
- Outcomes are captured to support reporting, learning, and continuous improvement.

7. Business Process

1. Request Lead submits a staffing request using a conversational front end (request intake)
2. The solution validates the request and stores it in the central repository.
3. AI identifies required capabilities and matches them against delivery, experience, interests, and capacity data.
4. The solution recommends a Pod Lead, supporting Pod Members, and optional alternates.
5. The Request Lead reviews the recommendation and adjusts it if needed.
6. The approved assignment is communicated to the Pod Lead.
7. The Pod Lead mobilizes the team and executes the request.

8. Outcome, status, and reporting data are captured at closeout.

Main Components of Future State Solution:

- 1. Request Management**

- Central request repository
- Standard intake Prioritization

- 2. Resource Management**

- Experience Profiles
- Capacity Profiles Interest Management

- 3. AI**

- Recommend Pod Leads
- Recommend Specialists/Pod Members
- Explain recommendations
- Learn from outcomes

- 4. Governance**

- Human approval
- Audit trail
- Role-based security

- 5. Reporting**

- Open requests
- Capacity
- Staffing metrics
- Portfolio dashboard

8. Functional Requirements (FR-001 ... FR-0XX)

The following functional requirements describe what the solution must do. They are written to be testable and traceable to the business objectives.

ID	Functional Requirement
FR-001	The system will allow an authorized Request Lead to create a staffing request .
FR-002	The system will validate mandatory request fields before a request can be submitted.
FR-003	The system will store each request in a central repository.
FR-004	The system will maintain a delivery taxonomy that describes the type of work requested.
FR-005	The system will maintain a capability taxonomy that maps work to required capabilities.
FR-006	The system will maintain experience profiles for Pod Leads and Pod Members.
FR-007	The system will maintain interest data as part of the resource profile.
FR-008	The system will maintain capacity data that reflects availability and active pod commitments.
FR-009	The system will generate an AI recommendation for Pod Lead and Pod Members based on request data and resource profiles.
FR-010	The system will provide a rationale for each AI recommendation.
FR-011	The system will allow a Request Lead to review and override a recommendation prior to approval.
FR-012	The system will require human approval before a staffing assignment becomes active.
FR-013	The system will notify the Pod Lead once the assignment is approved.
FR-014	The system will allow a Pod Lead to confirm execution readiness after assignment.
FR-015	The system will capture request status changes from submission through completion.
FR-016	The system will store the assignment history for reporting purposes.
FR-017	The system will capture outcome data and closeout information for completed requests.
FR-018	The system will provide dashboards for open requests, capacity, staffing metrics, and portfolio views.
FR-019	The system will support role-based access to requests, profiles, recommendations, and reports.
FR-020	The system will retain an audit trail of request creation, recommendation changes, approvals, and closeout actions.

9. AI Recommendation Framework

The AI recommendation capability is advisory. It does not replace human decision-making.

Inputs

- Request description and business context.
- Delivery taxonomy and required capabilities.
- Experience history and prior deliverables.
- Interests and development preferences.
- Capacity and active pod commitments.
- Historical outcomes, where available.

Outputs

- Recommended Pod Lead.
- Recommended Pod Members / specialists.
- Alternate candidates, if applicable.
- A reasoned explanation of why the recommendation was made.

Operating Rules

- The recommendation must be reviewable by a human before assignment.
- The system must support override and adjustment by the Request Lead.
- The system should preserve the evidence used to generate each recommendation.
- The system should support learning from outcomes over time, subject to governance approval.

10. Data Model & Taxonomies

The MVP data model should include the following core entities and taxonomies.

Entity	Description
Request	Stores request metadata, priorities, dates, status, and ownership.
Delivery Taxonomy	Defines the types of work (the types of requests) being requested and that needs to be delivered.
Role/Capability Taxonomy	Defines the capabilities required to deliver each work type or request.
Experience Profile	Tracks prior deliverables, role history, domain expertise, and learning stage.
Interest Profile	Captures areas of interest or development preference.
Capacity Profile	Tracks availability, allocation, and active pods.
AI Recommendation	Stores the recommendation, rationale, and supporting inputs.
Pod Assignment	Stores approved assignments and pod composition.
Outcome / Closeout	Stores completion status, lessons learned, and reporting data.

11. User Personas & Roles

Role	Primary Responsibilities
Request Lead	Creates requests, reviews recommendations, approves or adjusts staffing decisions, and communicates approved requests.
Pod Lead	• Receives approved requests, mobilizes the pod, leads execution, and maintains their own profile information (updates experience and interest data)
Pod Member / Specialist	Participates on pods, updates experience and interest data, and contributes to delivery.
Governance Lead	Reviews conflicts, monitors capacity, and supports portfolio review and governance.
Leadership	Reviews dashboards, approves priorities, and measures outcomes.
System Administrator	Manages configuration, roles, access, and operational support.

12. Business Rules

- Every request must have a human owner.
- AI recommendations are advisory and never final without approval.
- A Pod Lead must be identified for each approved request
- Capacity must be reviewed before assignment finalization.
- Profiles should be maintained by the relevant resource owner and/or governance process.
- Audit data must be retained for staffing actions, overrides, and approvals.

13. Reporting & Analytics

The solution should provide reporting at the operational and leadership levels.

Report / Dashboard	Key Measures
Open Requests	Count, time request is open, priority, request type, and ownership.
Capacity View	Availability, allocation, active pods, and constrained resources.
Staffing Metrics	Time to staff, approval lag, acceptance rate, and override rate.
Portfolio Dashboard	Demand by work type, trend, bottlenecks, and delivery status.
AI Effectiveness	Recommendation acceptance, override reasons, and outcome quality.

14. Governance & Security

- Role-based access control will be used.
- Access to requests, profiles, and recommendations will be limited to authorized users.
- Every recommendation review, approval, and override will be auditable.
- The solution should align with enterprise identity and authentication standards.
- Governance should define ownership of taxonomies, profiles, and capacity data.

15. Non-Functional Requirements

NFR Category	Requirement
Performance	Recommendation results should be returned within an acceptable user response time for standard requests.
Availability	The MVP should support the expected business operating window for request handling and reporting.
Scalability	The solution should support increased request volume and broader adoption after the pilot.
Security	The solution will protect request, profile, and recommendation data using role-based controls.
Usability	The request intake and review experience should be intuitive for business users.
Maintainability	Taxonomies, rules, and dashboards should be maintainable without extensive rework.
Auditability	All material actions must be traceable to a user, time, and event.

16. User Stories & Acceptance Criteria

- As a Request Lead, I want to submit a staffing request so that work can be staffed consistently.
 - Acceptance: a request can be created, validated, saved, and routed for recommendation.
- As a Request Lead, I want to review AI recommendations so that I can make an informed staffing decision.
 - Acceptance: the recommendation, rationale, and alternates are visible before approval.
- As a Pod Lead, I want to see my approved assignments so that I can mobilize the team quickly.
 - Acceptance: approved assignments appear in my queue and trigger notification.
- As a Governance Lead, I want to review exceptions so that capacity and policy risks are managed.
 - Acceptance: exceptions are logged and accessible for review.
- As Leadership, I want dashboards so that I can monitor staffing effectiveness and demand.
 - Acceptance: key metrics are available by dashboard and reflect current data.

17. MVP Scope

- Request intake and central request repository.
- Core delivery and capability taxonomies.
- Experience and capacity profile capture.
- AI-assisted recommendation for Pod Lead and Pod Members.
- Human review and approval workflow.
- Basic dashboards and audit trail.
- Pilot with one team or one operating group.

18. Roadmap

Phase	Description
Phase 1 – MVP Build	Stand up the core repository, taxonomies, profiles, recommendation logic, and reporting.
Phase 2 – Pilot	Validate the process with one team and refine the workflow based on feedback.
Phase 3 – Stabilization	Improve usability, reporting, and recommendation quality; resolve pilot findings.
Phase 4 – Scale	Extend adoption to additional teams and broaden governance and reporting.

19. Risks, Assumptions & Dependencies

Risks

- Incomplete or outdated profile data may reduce recommendation quality.
- Unclear ownership of taxonomies may create governance delays.
- Users may continue using informal channels if intake adoption is low.
- Scope creep could delay the MVP if too many features are introduced too early.

Assumptions

- The organization will support a standardized intake process.
- Users will maintain experience, interest, and capacity data with appropriate cadence.
- Human approval remains mandatory for staffing decisions.
- The pilot team will provide timely feedback.

Dependencies

- Agreement on taxonomy definitions and ownership.
- Access to relevant user, profile, and operational data.
- Alignment on governance and approval workflow.
- Availability of development and pilot stakeholders.

20. Success Metrics

Metric	Target / Expectation
Average staffing time	Target reduction versus current baseline.
AI recommendation acceptance rate	High percentage of recommendations accepted with minimal override.
Request visibility	All requests tracked in the central repository.
Utilization visibility	Leaders can see capacity and workload trends.
Pilot satisfaction	Positive feedback from the pilot team and leadership.

Appendix A. Glossary

Term	Definition
Request Lead	Person responsible for creating and managing the staffing request.
Pod Lead	Person accountable for mobilizing and leading the assigned Pod.
Pod Member	Contributor assigned to the Pod to execute work.
Capacity	Availability of a person to take on new work.
Experience Profile	Structured record of completed work and relevant history.
Delivery Taxonomy	Structured list of work types delivered by the organization.
Interest Taxonomy	Structured list of work types staff has experience in or would like to gain experience in.
Capability Taxonomy	Structured list of capabilities needed to deliver work types.

Appendix B. Traceability Matrix

Objective	Supporting Requirements
Business Objective: Reduce staffing time	FR-001, FR-003, FR-009, FR-012, FR-013, FR-018
Business Objective: Improve consistency	FR-004, FR-005, FR-009, FR-010, FR-020
Business Objective: Increase visibility	FR-003, FR-008, FR-015, FR-017, FR-018
Business Objective: Preserve governance	FR-011, FR-012, FR-019, FR-020
Business Objective: Scale the model	FR-004, FR-005, FR-006, FR-007, FR-015, FR-018

Appendix C. Mockups (Textual)

- Request Intake Screen: request details, taxonomy selection, required capabilities, and submit action.
- AI Recommendation Screen: recommended Pod Lead, Pod Members, rationale, alternates, and override controls.
- Approval Screen: review summary, decision, comments, and audit trail.
- Leadership Dashboard: open requests, staffing metrics, capacity view, and portfolio trends.

Appendix D. Diagrams (Textual) – to be provided at a later date

- Current-state vs future-state process flow.
- AI recommendation workflow.
- Data model / entity relationship diagram.
- Governance and approval flow.
- Dashboard layout overview.

Open Questions for Discovery and Discussion

- Should the MVP leverage an existing Oracle platform or use a lightweight custom solution?
 - Exploration completed with 2 teams; there are no existing solutions that can be leveraged
 - Aligned to create a simple/lightweight custom solution
- What should be the system of record for skills, experience, and capacity?
- How should employee profiles be maintained?
- What level of automation is appropriate for the MVP?
- What integrations are required?
- How will AI recommendations be evaluated during the pilot?