

MEETING SCRIPT

AI Pod Staffing

Frontend Prototype Walkthrough

A clear, word-for-word guide for presenting the current application

Recommended duration: 12-15 minutes

Start persona: Operations Lead | Start page: Command Center

1. Opening and meeting framing

Set expectations before showing any screen.

ON SCREEN

Select Operations Lead and open Command Center. Close any drawers or chat panels.

SAY

Today I will walk through the current AI Pod Staffing frontend prototype. The objective is to bring staffing demand, customer-defined capabilities, employee capacity, availability and AI-assisted recommendations into one governed workflow. The prototype is backed by the staffing workbook and uses the customer's project, deliverable and capability terminology. AI recommendations remain advisory, and a person must review the proposed pod before any staffing decision is final.

What the prototype demonstrates

- Standardized staffing-request intake.
- Customer-controlled project, deliverable and capability terminology.
- Role-scoped views for Operations Lead, Pod Lead and POD Member.
- Evidence-based fitment recommendations with human approval.
- Capacity, availability, governance and reporting experiences.

IMPORTANT

Call this a frontend and workflow prototype. Do not describe it as a completed production backend or a deployed autonomous OCI agent.

Simple workflow to remember

1	2	3	4	5
Capture request	Match capabilities	Check capacity	Recommend pod	Human approves

2. Operations Lead Command Center

Explain the four KPIs as one connected business story.

ON SCREEN

Point to the four KPI cards, then the staffing queue, Capacity Watch, Upcoming Demand and Audit Trail.

SAY

The Command Center gives the Operations Lead a concise view of staffing demand, team capacity, fulfillment progress and recommendations awaiting review. It helps the user understand what needs attention without placing detailed reporting on the homepage.

KPI	What it means	How to explain it
Open requests	Active requests that are not closed.	The badge identifies how many are high priority.
Team allocation	Average allocation across people visible to this persona.	Constrained means the person is at or above the prototype's 70% warning level.
Staffing progress	Staffed active requests divided by all active requests.	For example, 1 of 3 staffed equals 33%. The value changes with the live request list.
Pending recommendations	Candidate recommendations awaiting a recorded human decision.	Advisory means AI has suggested candidates but has not assigned anyone.

SAY

The Priority Staffing Queue shows requests that need recommendation, review or approval. Capacity Watch identifies workload pressure. Upcoming Demand shows when current requests are expected to require staffing. The Audit Trail card provides access to governance evidence without filling the homepage with operational detail.

KPI CAUTION

Read the live values from the screen. A locally created request may change Open Requests and Staffing Progress during the same browser session.

3. Staffing Requests and request creation

Show how demand is captured in customer language.

ON SCREEN

Open Requests. Show search, priority, project, deliverable and status filters. Open the status list and point out Closed.

SAY

This workspace is the central register of staffing requests. Each request is expressed using customer terminology: project type, key deliverables, required capabilities, request source, status and needed-by date. Users can search, filter and export the requests within their access scope.

ON SCREEN

Click New Request. Do not save during the main demonstration.

Explain the form in this order

- Request Title - a concise name for the work.
- Project Type - selects the relevant customer project category.
- Request Source - identifies the person or team asking for the work.
- Project Description - provides context beyond the title.
- Key Deliverables - supports multiple mapped deliverables and a request-specific Other value.
- Needed By Date, Estimated Start Date and Estimated Completion Date - separates urgency from delivery planning.
- Estimated Effort - entered in days, weeks or months.
- Requested Pod Size - defines the expected lead and contributor structure.
- Required Capabilities - defaults from selected deliverables but supports additions, removals and custom entries.
- Business Objectives and Expected Outcomes - captures audience, success measures and intended outcome.

SAY

The request form creates the business and delivery context needed to evaluate a staffing recommendation. The customer mapping supplies a controlled starting point, while request-specific changes remain possible.

72 HOURS

For existing workbook records, 72 hours is estimated effort, not elapsed completion time. New requests can capture effort in days, weeks or months.

4. AI Fitment Review

Explain what fitment means and how the reviewer changes requests.

ON SCREEN

Open AI Fitment. Use the new Staffing Request dropdown and select REQ-1042 - AI service launch communication.

SAY

AI Fitment means matching a staffing request with the most suitable Pod Lead and contributors. It is not an employee performance score. It answers: based on the required work, capabilities, experience, capacity and availability, who is the strongest fit for this request?

Explain the screen from left to right

- Request summary - dates, deliverables, effort, pod size, required capabilities, status and business outcomes.
- Staffing Request dropdown - switches the entire review between requests visible to the selected persona.
- Recommendation factors - the intended evidence categories used to explain the match.
- Recommended Pod Lead - the proposed person for overall coordination.
- Recommended Contributors - people proposed to complete the capability coverage of the pod.
- Rationale - a business-readable reason for each recommendation.

Factor	Prototype weighting	Meaning
Interest strength	35%	Strength of the person's relevant capability evidence.
Relevant delivery history	25%	Evidence from similar completed work.
Available capacity	25%	Current workload and recorded availability.
Growth preference	10%	Alignment with development interests.
Team continuity	5%	Benefit of maintaining an existing team relationship.

FITMENT SCORE

A score such as 94 represents suitability for this specific request. It is not a 94% employee-performance rating and it is not a guarantee of delivery success.

ON SCREEN

Select REQ-1041 briefly to show that a request without a recorded recommendation displays No AI fitment available. Return to REQ-1042.

5. Human approval and Agent Execution

Connect transparency with the mandatory human checkpoint.

ON SCREEN

On AI Fitment, click Approve Pod, explain the checkpoint, then click Cancel.

SAY

The AI does not make the final assignment. The reviewer must approve, adjust or decline the proposed pod. The production workflow will record the decision, reason, selected people and evidence version for audit.

ON SCREEN

Open Agent Execution, select REQ-1042 and click Run Fitment.

Execution stages

- Validate the staffing request.
- Resolve mapped and request-specific capabilities.
- Retrieve people and their supporting evidence.
- Apply availability, allocation and conflict checks.
- Rank candidates and explain the evidence.
- Pause for human review.

SAY

The purpose of this screen is transparency. The user should be able to see which evidence was considered, which constraints were applied and why the result was produced. In the current prototype, this flow visualizes the intended orchestration and reads the recommendation evidence stored in the workbook.

DO NOT OVERCLAIM

Agent Execution is currently a workflow visualization. It is not yet a deployed OCI or AIDP agent performing live enterprise scoring.

6. Allocation, Team & Skills and My Availability

Show how operational evidence supports better recommendations.

ON SCREEN

Open Allocation Calendar. Show capability and capacity filters, Show Alternates and the weekly grid.

SAY

The Allocation Calendar brings assignments, workload, travel and availability together. It allows the scheduler to inspect the week before confirming work and supports alternate suggestions when the preferred person is constrained.

ON SCREEN

Click Allocate, show the guardrail notice, then cancel.

SAY

Before saving an allocation, the intended guardrail checks whether the person would exceed the configured load limit. If the limit would be exceeded, the allocation is blocked and an alternate is proposed for human review.

ON SCREEN

Open Team & Skills. Point to Alex Rivera's profile and the team directory.

SAY

Team & Skills contains the resource evidence used by staffing recommendations: mapped capabilities, strength, supporting experience, current allocation and active pods. The capability names are loaded from the customer-controlled Skills or Type of Work taxonomy.

ON SCREEN

Open My Availability.

SAY

Employees can record leave, travel, training, reduced hours or other capacity changes before recommendations are generated. The prototype displays recorded facts and does not invent future capacity when supporting data is unavailable.

7. Persona-based experience

Demonstrate least-privilege views without spending too long on every persona.

ON SCREEN	Return to Command Center. Switch from Operations Lead to Pod Lead.	
SAY		
The Pod Lead sees only assigned requests and the people related to those requests. Allocation Calendar, Agent Execution, Reports and Administration are blocked for this persona.		
ON SCREEN	Click one locked workspace to show the access-blocked message. Then switch to POD Member.	
SAY		
The POD Member receives a personal experience. The general staffing queue is replaced by My AI Fitment Preview, and Capacity Watch becomes My Capacity. Requests, Team & Skills and My Availability are restricted to the signed-in member's own information. Allocation Calendar, Agent Execution, Reports and Administration remain blocked.		
Persona	Visible scope	Blocked workspaces
Operations Lead	Full prototype operational scope	None in the current prototype
Pod Lead	Assigned requests and related pod members	Allocation Calendar, Agent Execution, Reports, Administration
POD Member	Own requests, personal fitment, own profile and availability	Allocation Calendar, Agent Execution, Reports, Administration
DEMO CAUTION	Do not open the full POD Member AI Fitment page during the main walkthrough. Re-run and Approve controls still require persona-level UI refinement. Switch back to Operations Lead after the role demonstration.	

8. Reports, Administration and Ask AI Pod

Finish the tour with governance, insight and conversational access.

ON SCREEN

Switch back to Operations Lead and open Reports.

SAY

Reports contains the operational detail intentionally removed from the homepage. It shows demand, average allocation, staffing progress, pending recommendations, project types, deliverables, customer capabilities, people and current recommendation outcomes. The figures are described as a current workbook snapshot; the prototype does not manufacture historical trends.

ON SCREEN

Open Administration and move through Roles & Access, Taxonomies, Agent Rules and Audit Trail.

SAY

Administration demonstrates the governance layer. Roles & Access presents the intended permission model. Taxonomies display the customer-controlled project, deliverable and capability catalogue. Agent Rules show load protection and the requirement for human approval. Audit Trail shows recorded recommendation evidence and decision status.

ON SCREEN

Open Ask AI Pod and choose Who has capacity? or Explain REQ-1042.

SAY

Ask AI Pod provides a conversational way to query the staffing workbook. Its answers respect the selected persona, so a POD Member cannot use the assistant to retrieve broader team information.

CURRENT BEHAVIOR

The assistant currently returns workbook-backed prototype answers through the application API. It is not yet the final enterprise agent implementation.

9. Feedback completion, likely questions and close

End confidently while being accurate about what remains prototype-only.

Feedback now represented in the frontend

- Governance Lead renamed to Operations Lead.
- Employee renamed to POD Member.
- View Allocation removed from the homepage.
- Homepage reporting subtext removed; reporting retained in Reports.
- Closed added to the request status filter.
- Create Request enhanced with descriptions, outcomes, dates, effort units, pod size, multiple/custom deliverables and editable capabilities.
- Pod Lead and POD Member views scoped; restricted workspaces blocked.
- POD Member dashboard changed to My AI Fitment Preview and My Capacity.
- Operations Lead governance activity reduced to an Audit Trail access card.
- AI Fitment now has a visible Staffing Request selector.
- Customer Deliverables KPI replaced with Staffing Progress.

Likely questions

Question	Short answer
Does saving update Excel?	No. Prototype actions remain local and the source workbook is unchanged.
Is AI assigning people automatically?	No. It recommends and explains; a human approves, adjusts or declines.
Is the score final?	No. Current scores and rationale are demonstration evidence from the workbook.
Where does the data come from?	data/ai-pod-staffing-prototype.xlsx, currently mapping version v260810.1.
Is role security production-ready?	No. The frontend demonstrates scoping; production requires API and database enforcement.
Where are pod status updates?	This workflow remains an open functional requirement.

SAY

To summarize, the prototype demonstrates the intended end-to-end experience. A request is captured using customer terminology. Deliverables provide default required capabilities. People are evaluated using skills, evidence, capacity and availability. Recommendations are explained, capacity guardrails are applied, and a human remains responsible for the final staffing decision. The next phase is to replace prototype interactions with persistent workflows, enterprise integrations, backend authorization, governed agent orchestration and auditable production data.

FINAL REMINDER

Keep the discussion focused on the business workflow and user experience. Use technical architecture only when someone specifically asks how the production solution will be implemented.